

sheet04_solutions

November 29, 2020

##Submitted by: **Group Name:**MI 4

Group Members: Anuja Negi, Debapratim Jana, Tobiasz Budzynski, Muthukumar Pandaram

```
[1]: import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

$$E(w_{t+1}) = E(w_t) + \nabla E \cdot (-\eta_t d_t) + \frac{1}{2}(-\eta_t d_t)^T H(-\eta_t d_t) + \frac{1}{3!}E^{(3)}(w_t)(-\eta_t d_t) \dots,$$

where the fourth and next terms would need to be expanded with Frechet Derivative and cubic and higher order form. \

Fortunately, we are interested only in first three terms of the Taylor expansion.

$$\Delta E^T \leq 0$$

$$\Delta E^T \approx -E(w_t) + E(w_t) + \nabla E \cdot (-\eta_t d_t) + (-\eta_t d_t)^T H(-\eta_t d_t)$$

Using inequality above, taking the scalar η in front of each term and dividing by η , we get (the product with Hessian is non zero for non-zero d , because H is positive definite)

$$\eta_t \leq \frac{\nabla E d_t}{d_t^T H d_t}$$

As E^T has a minimum at an unknown w^* , it's gradient is 0 there. Taylor approximation is equal to

$$E_{w_{t+1}}^T \approx E_{[w_t]}^T + \frac{1}{2}(w_{t+1} - w_t)^T H(w_{t+1} - w_t)$$

$$\nabla E_{[w_{t+1}]}^T = \nabla E_{[w_t]}^T + \frac{1}{2} * 2H(-\eta_t d_t) \quad (@)$$

If we want in next step get as close as possible to w^* , we shall assume that left side of equation is 0. Putting symbol $*$ over η .

$$\eta_t^* H(-d_t) = d_t.$$

To calculate $\eta \in \mathbb{R}$ from this vector equation, we can take a dot product with d_t (or with $\nabla E_{w_t}^T$) from left side, to get real numbers. We can use the equality $d^T d = ||d||^2$.

$$-\eta_t^* (\nabla E_{w_t}^T)^T H(d_t) = \nabla E_{w_t}^T \cdot d_t$$

Now, there is no point of taking d_t perpendicular to gradient (don't go up!), hence the Hessian thing is non-zero (positive definite, bilinear form) and we can divide.

$$\eta_t^* = -\frac{\nabla E_{w_t}^T \cdot d_t}{(\nabla E_{w_t}^T)^T H(d_t)}$$

Let's check! If such η^* really minimize average cost function E , the orthogonality of direction d with gradient at the new point w_{t+1} should hold. It would mean, that level of E is tangent to direction d . So w_{t+1} is critical. Knowing, that E is convex we would be sure of minimality.

$$d_t \cdot \nabla E_w^T|_{w_{t+1}} \stackrel{?}{=} 0$$

Using definitions of η^* , and if direction is gradient $d_t = -\nabla E_{w_t}^T = -H w_t$ (what if it is not gradient??) and equation (@), we have

$$\begin{aligned} d_t \cdot \nabla E^T|_{w_{t+1}} &= d_t \cdot H w_{t+1} = \\ &= d_t \cdot H(w_t + \eta^* d_t) = d_t \nabla E_{w_t}^T + \eta^* d_t^T H d_t = d_t \cdot d_t - ||d_t||^2 = 0 \end{aligned}$$

Part (a) : Gradient Descent

[2] : *# Initialise training data and weight matrix*

```
X = np.array([ [1 , -1] , [1, 0.3] , [1 , 2] ]).T
Y = np.array([-0.1, 0.5, 0.5])

w_init = np.array([-0.45, 0.2]).T
```

```
[3]: # (a) Gradient Descent

def gradient_descent(X, Y, w_init, max_iter, lr = 0.001):

    w= w_init.copy()

    mse_history=[]
    stopping_criteria=True
    w0_history = []
    w1_history = []
    iter = 0

    H = np.dot(X, X.T)

    while iter < max_iter and stopping_criteria:
        g = np.dot(H,w) - np.dot(X, Y.T)
        w = w - lr*g

        Y_pred = np.dot(w, X)
        mse=np.square(np.subtract(Y, Y_pred)).mean()
        # print(mse)
        mse_history+= [mse]

        #stopping criterion
        if iter>0:
            mse_new=mse_history[iter]
            mse_old=mse_history[iter-1]
            if ((abs(mse_new-mse_old))/mse_old)<=0.005:
                stopping_criteria=False

        w0_history += [w[0]]
        w1_history += [w[1]]
        iter += 1

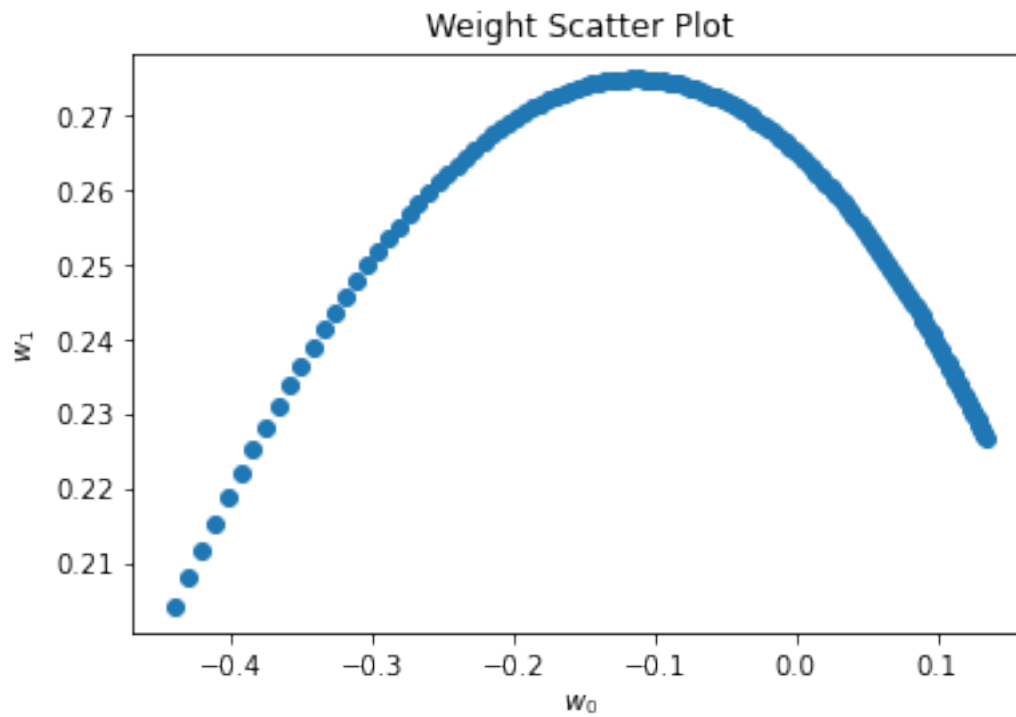
    return w0_history , w1_history, mse_history
```


```
[4]: w0_history , w1_history, mse_history = gradient_descent(X, Y, w_init, max_iter=
    ↳ 1_000, lr = 0.005)

plt.scatter(w0_history , w1_history)
plt.xlabel('$w_0$')
plt.ylabel('$w_1$')
plt.title('Weight Scatter Plot')
```

```
plt.show()

plt.plot(w0_history , label = '$w_0$')
plt.plot(w1_history , label = '$w_1$')
plt.ylabel('Weight ($w_i$) ')
plt.xlabel('Iteration')
plt.title('Time Evolution of Weight Components')
lgd = plt.legend()
plt.show()
```






As we can see, the weight vector components eventually stabilise when some optima (minima) has been found in the parameter weight space.

```
[5]: w_final = np.array([w0_history[-1] , w1_history[-1] ]).T

Y_pred = np.dot(w_final, X)
plt.plot(Y, '.m' , label = 'Data')
plt.plot(Y_pred, 'xr', label = 'Predicted')
plt.title('Gradient Descent : Predicted vs Original Data')
plt.show()
```



```
[6]: # (b) Line search
def line_search(X, Y, w_init, max_iter, stopping = True):

    w= w_init.copy()

    w0_history = []
    w1_history = []
    mse_history=[]
    iter = 0
    stopping_criteria=True

    H = np.dot(X, X.T)

    while iter < max_iter and stopping and stopping_criteria:
        g = np.dot(H,w) - np.dot(X, Y.T)
        lr = np.dot(g.T,g)/np.dot(g.T, np.dot(H, g))
        w = w - np.multiply(lr, g, dtype=np.longdouble)

        Y_pred = np.dot(w, X)
        mse=np.square(np.subtract(Y, Y_pred)).mean()
        # print(mse)
        mse_history+= [mse]
```

```

    #stopping criterion
    if iter>0:
        mse_new=mse_history[iter]
        mse_old=mse_history[iter-1]
        if ((abs(mse_new-mse_old))/mse_old)<=0.005:
            stopping_criteria=False

    w0_history += [w[0]]
    w1_history += [w[1]]
    iter += 1

    return w0_history , w1_history, mse_history

```

```

[7]: w0_history , w1_history, mse_history = line_search(X, Y, w_init, max_iter = 50)

```

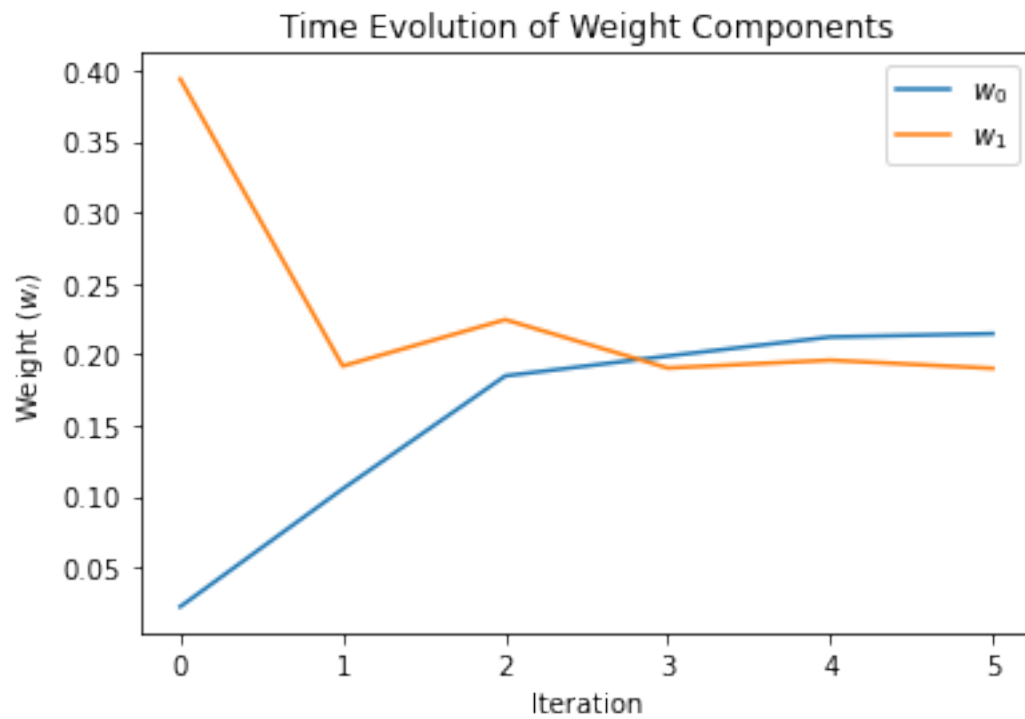

```

plt.scatter(w0_history , w1_history)
plt.xlabel('$w_0$')
plt.ylabel('$w_1$')
plt.title('Weight Scatter Plot')

plt.show()


plt.plot(w0_history , label = '$w_0$')
plt.plot(w1_history , label = '$w_1$')
plt.ylabel('Weight ($w_i$) ')
plt.xlabel('Iteration')
plt.title('Time Evolution of Weight Components')
plt.legend()
plt.show()

```




```
[8]: w_final = np.array([w0_history[-1] , w1_history[-1] ]).T

Y_pred = np.dot(w_final, X)
plt.plot(Y, '.m' , label = 'Data')
plt.plot(Y_pred.round(), 'xr', label = 'Predicted')
plt.title('Gradient Descent : Predicted vs Original Data')
plt.legend()
plt.show()
```



```
[9]: # (c) Conjugate Gradient
def conjugate_gradient(X, Y, w_init, max_iter, stopping = True):

    w= w_init.copy()

    mse_history=[]

    w0_history = []
    w1_history = []
    iter = 0
    stopping_criteria=True

    H = np.dot(X, X.T)
```

```

while iter < max_iter and stopping and stopping_criteria:
    if iter==0:
        g_old = np.dot(H,w) - np.dot(X, Y.T)
        d=-g_old
        lr = -np.dot(d.T,g_old)/np.dot(d.T, np.dot(H, d))
        w= w + lr*d
        g_new = np.dot(H,w) - np.dot(X, Y.T)
        beta=-np.dot(g_new.T,g_new)/np.dot(g_old.T,g_old)
        d=g_new+beta*d
        g_old=g_new
    else:
        lr =-np.dot(d.T,d)/np.dot(d.T, np.dot(H, d))
        w = w + lr*d
        g_new = np.dot(H,w) - np.dot(X, Y.T)
        beta=-np.dot(g_new.T,g_new)/np.dot(g_old.T,g_old)
        d=g_new+beta*d
        g_old=g_new

    Y_pred = np.dot(w, X)
    mse=np.square(np.subtract(Y, Y_pred)).mean()
    mse_history+= [mse]

    #stopping criterion
    if iter>0:
        mse_new=mse_history[iter]
        mse_old=mse_history[iter-1]
        if ((abs(mse_new-mse_old))/mse_old)<=0.005:
            stopping_criteria=False

    w0_history += [w[0]]
    w1_history += [w[1]]

    iter += 1

return w0_history , w1_history, mse_history

```

```

[10]: w0_history , w1_history, mse_history = conjugate_gradient(X, Y, w_init,
    ↪max_iter = 50)


plt.scatter(w0_history , w1_history)
plt.xlabel('$w_0$')
plt.ylabel('$w_1$')

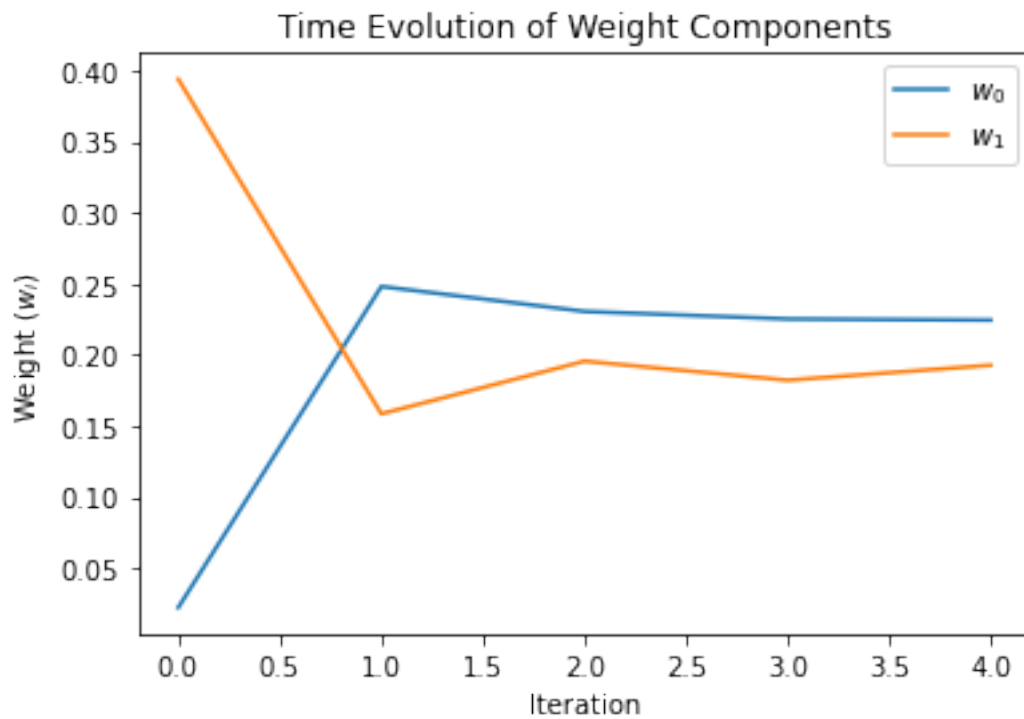
```

```
plt.title('Weight Scatter Plot')

plt.show()


plt.plot(w0_history , label = '$w_0$')
plt.plot(w1_history , label = '$w_1$')
plt.ylabel('Weight ($w_i$) ')
plt.xlabel('Iteration')
plt.title('Time Evolution of Weight Components')
plt.legend()
plt.show()
```






```
[11]: w_final = np.array([w0_history[-1] , w1_history[-1] ]).T

Y_pred = np.dot(w_final, X)
plt.plot(Y, '.m' , label = 'Data')
plt.plot(Y_pred.round(), 'xr', label = 'Predicted')
plt.title('Gradient Descent : Predicted vs Original Data')
plt.legend()
plt.show()
```



```
[12]: w0_history , w1_history, mse_history_gd = gradient_descent(X, Y, w_init,
    ↪max_iter = 1_000, lr = 0.005)
w0_history , w1_history, mse_history_ls = line_search(X, Y, w_init, max_iter =
    ↪50)
w0_history , w1_history, mse_history_cg = conjugate_gradient(X, Y, w_init,
    ↪max_iter = 50)

plt.plot(mse_history_gd,label='Gradient Descent')
plt.plot(mse_history_ls,label='Line Search')
plt.plot(mse_history_cg, label= 'Conjugate Gradient')
plt.ylabel('Mean squared error between the prediction and ground truth ')
plt.xlabel('Iteration')
plt.title('Convergence of errors for different gradient methods')
plt.legend()
plt.show()
```



It can be seen from the above plot that the line search and conjugate gradient methods help the error to converge in a very less number of iterations compared to the general gradient descent method.